

Access Paths

INFS602 Physical Database Design



Learning Outcomes

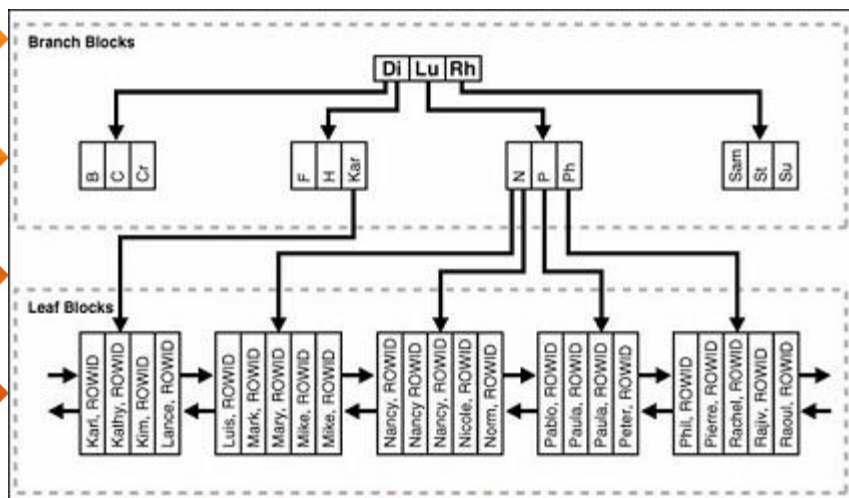
- Understand the role played by access paths in improving performance
- Understand the trade-offs between choosing *Indexing*, *Hashing* and *Clustering* in Physical Design
- Examine Oracle's Access Path mechanisms

Indexing



INDEX


ALCOHOL CAN BE A GAS!



- A. zwimert 94
 abalone, 294–295
 Abbey, Edward, 1f
 ABE, 439
 abiogenic oil, 61
 above-ground fuel tanks, 269, 269f
 absorption, **35**
 acceleration
 combustion and, 344
 engine conversion and, 328, 331, 529
 unmodified vehicle tests and, 331
 accelerator pump, **354**, 369–372
 acetaldehyde, **348**
 acetic acid, **109**
 acetone, **137**
 in ABE, 439–440
 acid(s). *See also* acidity; fatty acids
 bases and, 115–117
 emergencies with, 115–116
 sulfuric, 54, 115
 acid hydrolysis, **132**
 strong, 132–133
 weak, 132–133, 134, 136, 137
 acid rain, 54, 57–58
 acid/low-proof alcohol, 353, 354
 acidity, fermentation process and, **84**
 acridine, 55–56
 Acura Integra, 389, 529–532, 531f
 additives
 gasoline, 350, 356, 357–358, 360
 lubrication improvement, 425
 petroleum, 392
 Adelman, H.G., 454
 ADM. *See* Archer Daniels Midland
 adsorption, **226–229**
 advance, ignition systems, 404–405
 adventurous, **120**, 128f
 AGE-85, 337
 agitation
 cooking and, 246–249
 definition of, 246
 distillation, 251–256
 drive belt sheaves for, 248
 feedstock fermentation, 94, 99, 101, 105,
 249–251, 481
 agitators, **94**, 105, 112, 232, 234–235, 244,
 246–250
 chain drive, 489
 motor, 249, 254
 tanks and, 249f
 vaporized alcohol-fueled engine for, 421f
 agribusiness, 513
 agriculture. *See also* permaculture;
 polyculture
 community-supported (CSA), **503–505**
 integrated food/energy, 513
 organic, 46–50, 317
 protein and, 31–32
 U.S. corn, 27, 27f, 31–32, 31f, 39–40
 Agropyre oregana (mushroom), 314f
 Agrol, 18
 Agrol Company, 17, 18
 air conditioners
 cogenerators as, 445
 heat pumps compared to, 218, 219
 household cogenerated, 445
 ice block, 447
 air pollution, 34–35, 56
 catalytic converters and, 379
 coal and, 57–58
 exhaust, 425
 neat ethanol reducing, 350
 small engine, 421
 stoichiometric ratios and, 379–380
 two-stroke engines and, 425
 wood smoke, 224, 339
 aircraft
 alcohol-fueled, 17, 73, 73f, 337–338, 337f
 engines of, 336–339
 vaporized alcohol fuel and, 336, 338, 339
 airflow engine tuner, 377f, **387**
 fuel delivery increase by, 387
 air/fuel ratio, 388, 400, 410, 528
 calculating, 364
 carburetion and, 364, 367, 368–369
 enriching/adjusting, 339, 365, 367, 368–
 369, 383, 410, 424, 529
 gasoline, 364, 379–380, 389f
 metering jet determining, 364
 oxygen sensor/catalytic converters and,
 379–380
 range of, 529
 stoichiometric ratio as, **364**
 temperature and, 410
 utility engine, 424
 air-to-air heat exchangers, 302
 Alaska, 152
 pipeline, 57f
 Alberta, oil processing sites in, 52–53
 alcohol, **9**. *See also* alcohol fuel; alcohol
 production; alcohol production
 process; ethanol; ethanol v. gasoline;
 proof
 beverage, 206, 208, 469
 boiling point of water and, 185, 191
 combustion of, **35**, 446
 cooking food with, 339–341
 cooling with, 447
 denatured/denaturants for, 268, 270–271,
 274, 327, 394, 429, 500
 developing countries and, 41, 339–341,
 340f
 dry/drying, 225, 226–227, 227f, 236–237
 dual markets for, 71
 enrichment of, **192**, 192f, 194, 206
 flashpoint of, **268**, 269f
 forms of, 437
 generator using, 444
 household power use of, 446–448
 industrial-grade, 206
 leakage of, 268
 lighting with, 447
 liquid, 210
 off-road uses for, 196–197, 339–341, 444,
 462
 oxygen content of, 347
 phase separation of, 225–226
 prairie v. corn, 42
 proof requirement and, 196–197
 reforming, 431
 sources for, 119–180
 storage of, 232, 269–274, 268f, 271f
 sugar, 136
 vaporized, 66, 331f, 332–333, 418
 wood, **437**
 yield calculation for, 111–112
 alcohol dryer, **227f**
 alcohol fuel. *See also* blend(s); dual-fuel
 capability; engine conversion;
 ethanol; proof; specific blends
 aircraft using, 17, 73, 73f, 337–338, 337f
 air/fuel ratio for, 339, 365, 367, 368–369,
 383, 410, 424, 529
 alcohol pumps for, 430
 average latent heat of, 344
 blends of, 16, 70, 328, 356–357, 450–451
 Brazil export of, 41
 carburetors and, 363–375, 364f
 castor oil and, 426, 451, 452
 cogeneration using, 441–446
 diesel blended with, 450–451
 engines built for, 333–336
 fuel injection and, 379
 fuel needs supplied by, 27, 27f
 gasoline mixed with, 225–226, 375, 431
 legal definition of, 430, 500
 low-proof/acidic, 353, 354
 mad cow disease and, 316
 misage and, 263, 351–353
 myths about, 24–37
 octane in, 337, 435
 100%, 333–336
 peak power of, 336
 polyculture production of, 29, 39–40
 price of, factors in, 459, 463, 501
 redefining, 500
 refrigerators, 341
 safety of, 356
 sale of, 33
 temperature, 410
 vaporization of, 332, 333f, 341, 399, 409–
 410, 418, 425
 vehicle warranties and, 468–469
 water content of, 356

What is an Index?

- An index is a fast access path to data
- An index is a table data structure created to improve performance in some cases
- An index consists of two fields, the index key and a pointer which gives the address of the actual record
- The index is sorted on the index key
- For very large files, it is necessary to have a multi-level index (called B tree indexes) as the index itself may be large



Indexes

- Indexes are created on one or more columns of a table.
- After it is created, an index is
 - automatically maintained and used by Oracle.
- Changes to table data (such as adding new rows, updating rows, or deleting rows) are automatically incorporated into all relevant indexes with complete transparency to the users.

Index Properties

Two kinds of blocks:

- Branch Blocks - contain index data that points to lower-level index blocks

They store:

- The minimum key prefix needed to make a branching decision between two keys
- The pointer to the child block containing the key


- Leaf blocks - contain every indexed data and a corresponding rowid for the actual row

They store:


- The complete index key value for each table row
- ROWID pointer for each table row

All <Key, RowID> pairs are linked to their left and right siblings (forward and backward chain) . They are sorted by (key, ROWID).


Branch Blocks




Block 001 (root block)



Block 002




Block 003




Block 007

Oracle B-Tree Index Structure


Leaf Blocks




Block 008




Block 009



Block 010



Block 011



Block 099

Index Unique Scan

Steps in Index Unique Scans

- Start with the root block
- Search the block keys in sequential order for the key range that contains the required value
- If found, then follow the link from this key value range to the child block
- If no key range contains the required value from Step 2, then the row does not exist
- Repeat steps 2 through 4 if the child block is a branch block
- Search the leaf block for key equal to the value
- If key is found, then return the ROWID of the first key that equals the required value
- If key is not found, then the row does not exist

Index Unique Scan

- If searching for 224:
 - In the root block, 200...250 is the range that contains the required key value = 224
 - Follow the link to branch block 7 (200...250)
 - In this branch block, 221...228 is the range that contains the required key value = 224
 - Follow the link to leaf block 11 (221...228)
 - In this leaf block, search for key 224
 - Found 224, return (KEY, ROWID)

Index Range Scan

- If searching for ≥ 224 (*do this yourself*)

Indexes

- Consider a file with 1 million records. Suppose that record size is 200 bytes, and the block size is 2000 bytes. Also suppose that the size of the index key is 14 bytes and the pointer size is 6 bytes.
- These 10^6 entries would be stored in $10^6/10^2 = 10^4$ blocks (10,000)
- Thus 10^4 entries are needed at the next level of the index – these entries would be stored in $10^4/10^2$ blocks (= 100)
- At the next level of the index 100 entries are needed – these can be stored in 1 block – so no further levels are necessary
- Thus the no. of indexing levels = 3

Indexes

- The example illustrates the power of indexing – a random search on a 200MB file involved just 4 block accesses
- However, it must be noted that, while an index speeds up retrieval it does slow down updates
- Thus we only build indexes on columns which have a high read/write ratio
- Apart from update overhead, indexes also have a fairly substantial space overhead – up to 25% for each index built, depending on the Database System

Bitmap Indexes

- In a regular index a list of rowids is stored for each key corresponding to the rows with that key value.
- In a *bitmap index*, a bitmap for each key value is used instead of a list of rowids.
- Each bit in the bitmap corresponds to a possible rowid, and if the bit is set, it means that the row with the corresponding rowid contains the key value
- A mapping function converts the bit position to an actual rowid

Bitmap Indexes

Bitmap Index Example- CUSTOMER Table

CUSTOMER #	MARITAL_ STATUS	REGION	GENDER	INCOME_ LEVEL
101	single	east	male	bracket_1
102	married	central	female	bracket_4
103	married	west	female	bracket_2
104	divorced	west	male	bracket_4
105	single	central	female	bracket_2
106	married	central	female	bracket_3

Bitmap Indexes

REGION='east'	REGION='central'	REGION='west'
1	0	0
0	1	0
0	0	1
0	0	1
0	1	0
0	1	0

- An analyst investigating demographic trends of the company's customers might ask, "How many of our married customers live in the central or west regions?"


```
SELECT COUNT(*) FROM CUSTOMER WHERE MARITAL_STATUS = 'married' AND REGION IN ('central','west');
```

Bitmap Indexes

status = 'married'		region = 'central'		region = 'west'					
0		0		0		0		0	
1		1		0		1		1	
1	AND	0	OR	1	=	1	AND	1	=
0		0		1		0		1	
0		1		0		0		1	
1		1		0		1		1	

- The advantages of using bitmap indexes are greatest for low cardinality columns: that is, columns in which the number of distinct values is small compared to the number of rows in the table
- For example, on a table with one million rows, a column with 10,000 distinct values is a candidate for a bitmap index

Bitmap Index Size





BTree Indexes vs. Bitmap Indexes

- Use BTree indexes for columns with high cardinality (e.g. customer name, phone number, etc.)
- Use bitmap indexes on columns with low cardinality (typically columns whose values are repeated more than a hundred times)

BTree vs Bitmap Indexes

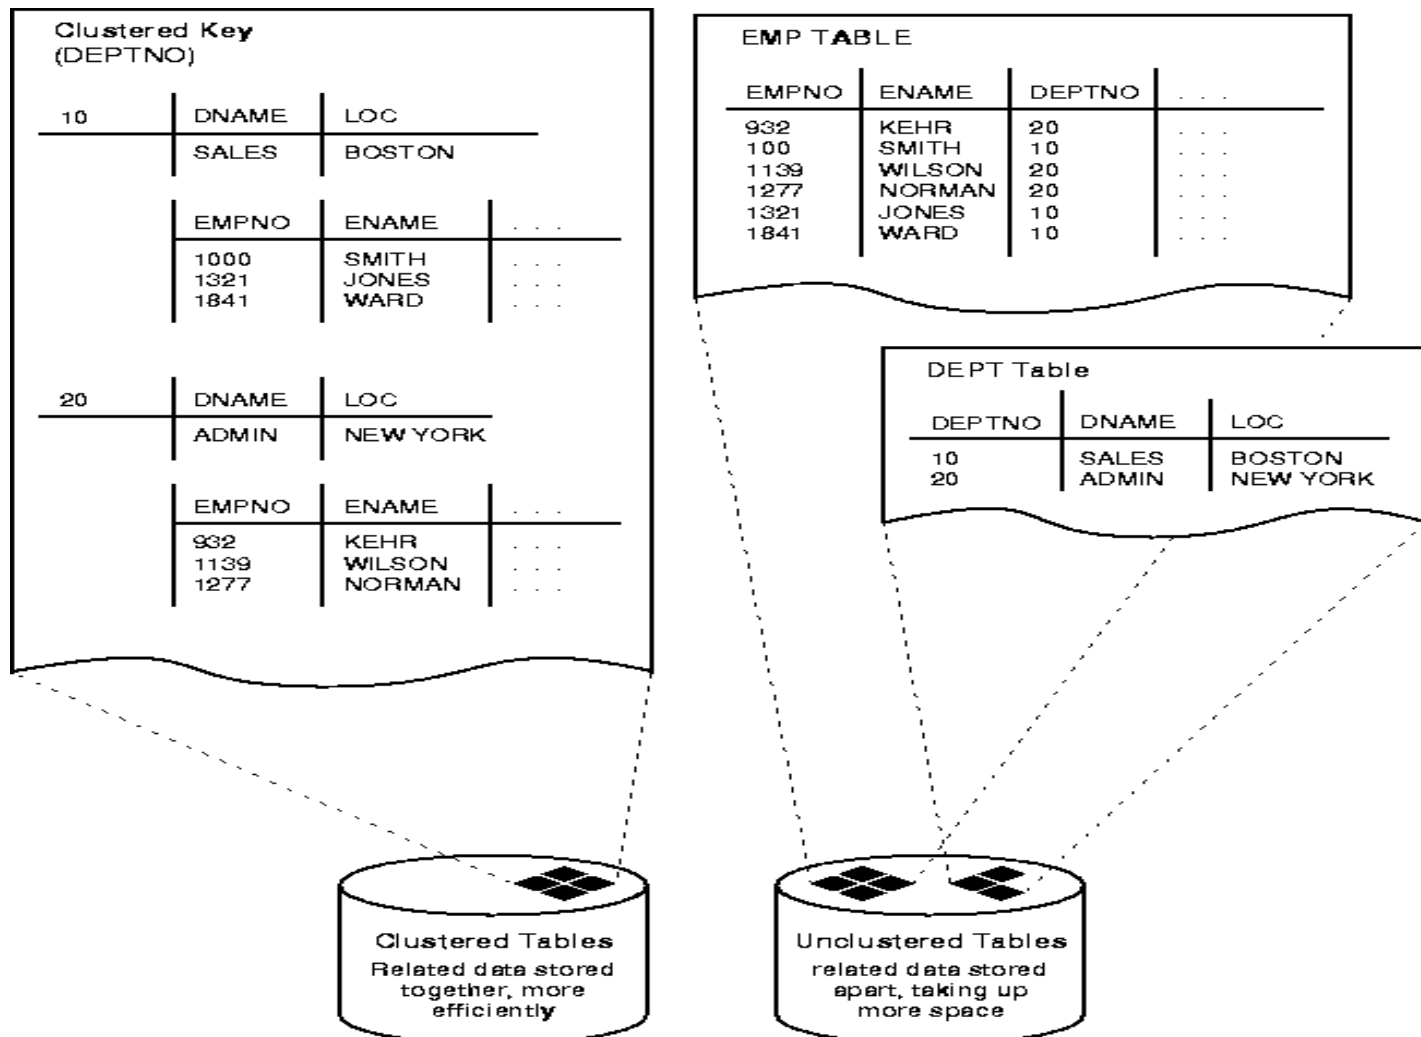
BTree	Bitmap
Suitable for high cardinality columns	Suitable for low cardinality columns
Update on keys relatively inexpensive	Updates on keys very expensive
Inefficient for queries using OR predicates	Efficient for queries using OR predicates
Useful for OLTP	Useful for data warehousing

BTree vs. Bitmap Indexes

- The following statements create indexes in Oracle:
 1. BTree `CREATE INDEX emp_ename_idx ON emp(ename)`
 2. Bitmap `CREATE bitmap INDEX emp_region_idx ON emp(region)`

Clusters

- Clusters store data from two or more related tables together in an interleaved fashion



Clustering

- In clusters rows from more than one table are stored in each data block
- Disk I/O is reduced and access time improves for joins of clustered tables (in a master-detail relationship)
- However tables that are *inserted* into or *modified* often are not good candidates for clustering
- In Oracle, clusters must be created before creating the tables that make up the cluster

```
CREATE CLUSTER emp_dept (deptno NUMBER(3)) PCTUSED 80 PCTFREE 5  
    SIZE 600)
```

- The cluster size specifies the average number of bytes required to store the rows for a cluster key value

Clustering

- Once the cluster is created the next step is to create the tables and the cluster index

```
CREATE TABLE dept ( deptno NUMBER(3) PRIMARY KEY, ... ) CLUSTER  
emp_dept (deptno)
```

```
CREATE TABLE emp ( empno NUMBER(5) PRIMARY KEY, ename  
VARCHAR2(15) NOT NULL, . . . deptno NUMBER(3) REFERENCES dept)  
CLUSTER emp_dept (deptno)
```

```
CREATE INDEX emp_dept_index ON CLUSTER emp_dept  
INITRANS 2 MAXTRANS 5 PCTFREE 5;
```


Hash Clustering

- Hashing goes one step better than indexing by calculating the locations of rows (rather than having to search using an index)
- Hashing provides optimal performance when the search criterion is an equality condition on a column value

e.g. Select all details of employee with empno =1234

Hash Clustering

- When a table is hashed its rows are stored in a number of *buckets*
- Each *bucket* contains rows which have the same hash key value
- A hashing *function* is used to map rows to *buckets*
- In order to find a row with only one block access (optimal value) the hashing function must distribute rows evenly across buckets

Hash Clustering

- The best hashing function is usually the remainder function used by the algorithm below

Algorithm

1. Determine the number of buckets to be allocated to the file
2. Select a prime number that is approx. equal to this number
3. Divide each hash key value by this number
4. Use the remainder as the bucket address

Hash Clustering Example

- Consider this Products table (15 products)

<u>Product #</u>	<u>Description</u>	<u>Finish</u>
100	Stereo Cabinet	Maple
125	Coffee Table	Walnut
153		
207		
219		
221		
286		
314		
363		
394		
418		
434		
488		
500		
515		

Hash Clustering Example

- If two rows fit into a bucket (block), then we need at least 8 buckets ($15/2 = 7.5$)
- Use a hash function based on 13 (prime):

Key	Prime Divisor	Product	Remainder (hash)	Bucket #	Key value
100	13	7	9	0	221
125	13	9	8		286
153	13	11	10	1	
207	13	15	12		
219	13	16	11	2	314
221	13	17	0		418
286	13	22	0	3	
314	13	24	2		
363	13	27	12	4	394
394	13	30	4		
418	13	32	2	5	434
434	13	33	5		
488	13	37	7	6	500
500	13	38	6		
515	13	39	8	7	488
				8	125
					515
				9	100
				10	153
				11	219
				12	207
					363

Hashing Performance Characteristics

- Hashing provides very good performance for searching, modifying and deleting rows when the search criterion is an equality condition on the hash key
- In order to ensure good performance, the number of overflows must be minimized, this is done by selecting the hash key carefully
- Hashing clusters the table, so only ONE hash key (which can be composite) can be used (contrast with indexing) for a given table
- Hashing does not support range searching at all (full table scan needed)

Oracle's Hashing Mechanism

- Oracle has a built-in hash function based on the “Remainder” method
- This function requires two parameters:
 1. Hashkeys – which specifies the total number of buckets required (Oracle will round this value up to the nearest prime number)
 2. Size – this gives the total space in bytes required to store the average number of rows associated with a bucket
- Oracle enables more than one table to be hashed in the same cluster

```
CREATE CLUSTER personnel (deptno NUMBER )
```

```
    SIZE 512 HASHKEYS 500 STORAGE (INITIAL 100K NEXT 50K);
```

```
CREATE TABLE emp (empno NUMBER PRIMARY KEY, ename  
    VARCHAR2(10) NOT NULL ,job VARCHAR2(9), hiredate DATE , sal  
    NUMBER(10,2), deptno NUMBER(2) NOT NULL )  
    CLUSTER personnel (deptno)
```

```
CREATE TABLE dept (deptno NUMBER(2), dname VARCHAR2(9), loc  
    VARCHAR2(9)) CLUSTER personnel (deptno);
```


Oracle's Hashing Mechanism

- To define a single table hash cluster use:

```
CREATE CLUSTER personnel (deptno NUMBER)
    SIZE 512 SINGLE TABLE HASHKEYS 500
```

- In cases when the cluster key is a unique identifier that is uniformly distributed over its range, you can by-pass Oracle's internal function and simply specify the hash column

```
CREATE CLUSTER emp_cluster (empno NUMBER)
    ...
    SIZE 55
    HASH IS empno HASHKEYS 10007
```

Oracle's Hashing Mechanism

- If the cluster key values are not evenly distributed you can create your own hash function
- For example, if your cluster key is the employee's home area code, it is likely that many employees will hash to the same hash value
- To alleviate this problem, you can place the following expression in the `HASH IS` clause of the `CREATE CLUSTER` command:

```
MOD (emp.home_area_code + emp.home_prefix + emp.home_suffix), 101)
```

Parallelism in Oracle

- Many DBMSs parallelize data-intensive operations to improve performance
- Improve performance through parallelize of various operations.
- parallel execution is useful for many types of operations that access significant amounts of data.
- Generally, parallel execution improves performance for:
 - Loading data (SELECT statements)
 - Queries.
 - Creation of large indexes.
 - Bulk INSERTs, UPDATEs, and DELETEs.
 - Aggregations and copying.

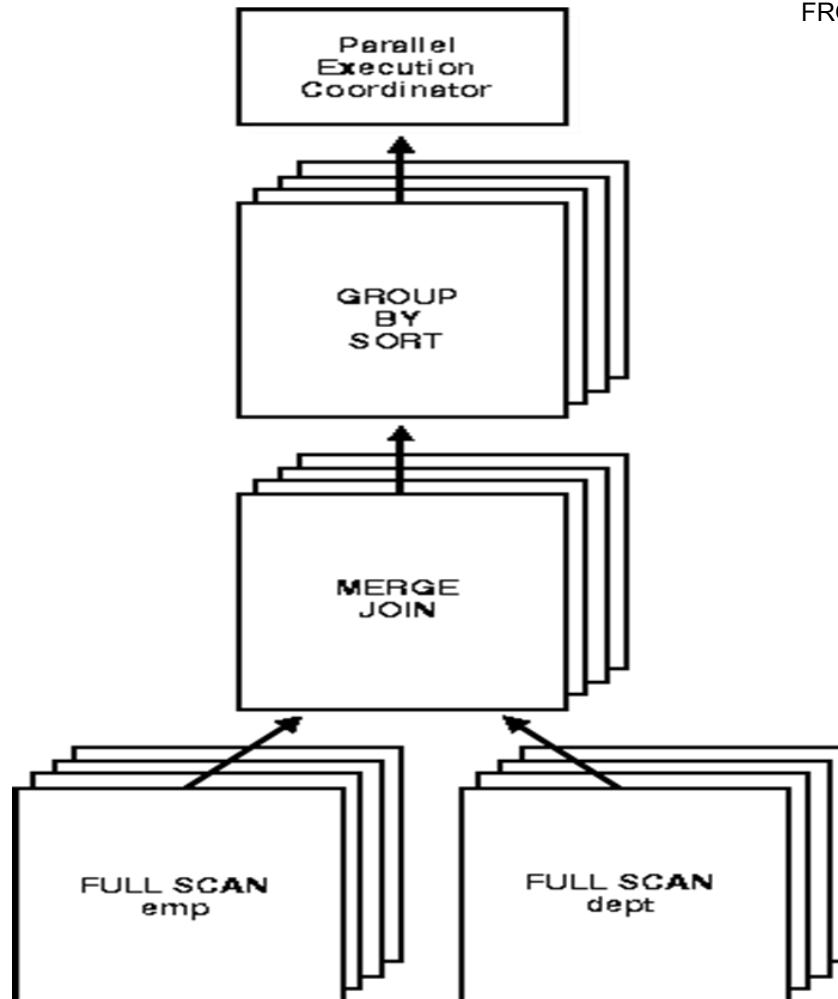
Using Parallelism to improve Performance

Two types of parallelism are possible:

1. Intra-operation parallelism – here a single operation is parallelised
2. Inter-operation parallelism

Using Parallelism

- Parallelism can also be used to perform many operations simultaneously



```
SELECT dname, MAX(sal), AVG(sal)
FROM emp, dept
      GROUP BY dname;
      WHERE emp.deptno = dept.deptno
```


Parallelism in Oracle

- Oracle supports both intra-operation and inter-operation parallelism
- Oracle allows the user to set the degree of parallelism
- The degree of parallelism is specified at the statement level (with hints or the PARALLEL clause), at the table or index level or by default based on the number of CPUs.

```
ALTER TABLE customer PARALLEL 4
```

```
ALTER CLUSTER dept_emp PARALLEL 4
```

```
ALTER INDEX iemp PARALLEL
```

```
SELECT /*+ PARALLEL(emp, 4) */ COUNT(*) FROM emp
```



Parallel Execution

- The user session or shadow process takes on the role of a coordinator, often called the query coordinator.
- The query coordinator obtains the necessary number of parallel servers.
- The SQL statement is executed as a sequence of operations (a full table scan to perform a join on a nonindexed column, an ORDER BY clause, and so on). The parallel execution servers perform each operation in parallel if possible.
- When the parallel servers are finished executing the statement, the query coordinator performs any portion of the work that cannot be executed in parallel. For example, a parallel query with a SUM() operation requires adding the individual subtotals calculated by each parallel server.
- Finally, the query coordinator returns any results to the user.

Creating a Summary Table in Parallel

Parallel Execution Coordinator

```
CREATE TABLE summary  
  (C1, AVGC2, SUMC3)  
  PARALLEL 5  
  AS  
  SELECT  
    C1, AVG(C2), SUM(C3)  
  FROM DAILY_SALES  
  GROUP BY (C1);
```





References

1. Fundamentals of Database Systems, Elmasri/Navathe, Chapters 5 and 6
2. Modern Database Management, Hoffer/Prescott/Topi, Chapter 6
3. Oracle 10g Concepts
4. Oracle 10g Administrator's Guide
5. Oracle 10g Tuning Guide

Review Questions

1. What is the decision to choose B* indexes over bitmap indexes based on?
2. Discuss the situations in which Hashing is preferred over Indexing.
3. How does clustering improve performance?
4. Discuss two different ways that parallelism can be used to improve performance in a large database environment