

Data Warehousing

INFS602 Physical Database Design




Agenda

- Data warehousing
- Operations
- Data warehousing strategies

What is a Data Warehouse?

- A Data Warehouse is a *subject-oriented, integrated, non-volatile, time-varying* repository of data
- It is a central location that data from different databases are stored and separated from operational database

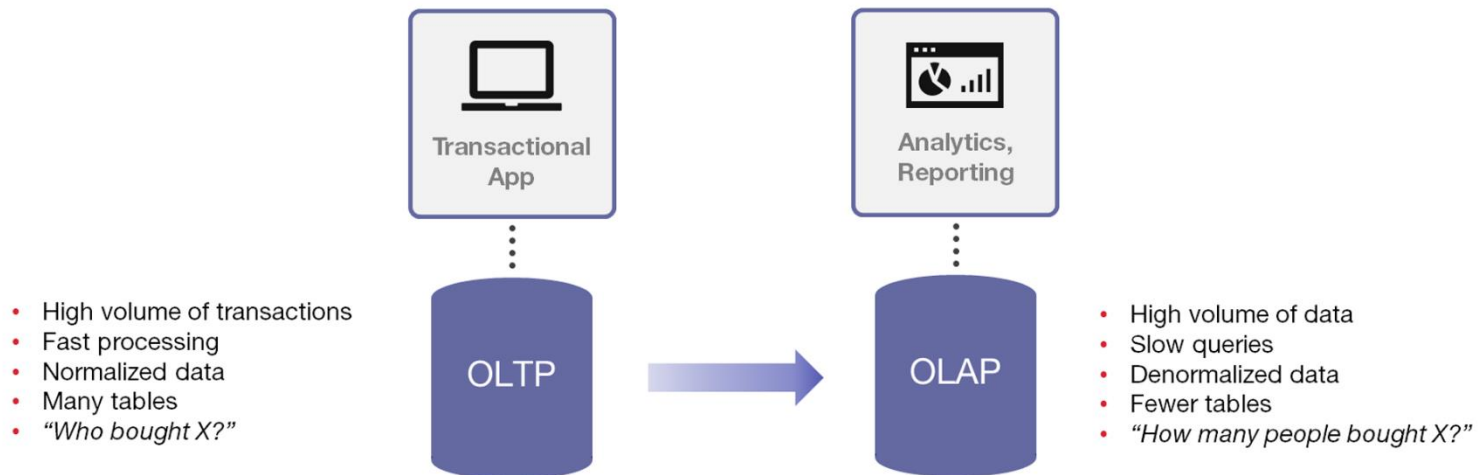


| Characteristic | OLTP Database | OLAP Data Warehousing |
|---------------------------------|---|---|
| Purpose | Supports transaction processing | Supports information requests |
| Source of data | Business transactions | Multiple files, databases—internal and external to firm |
| Data Access Allowed Users | Read and Write | Read Only |
| Primary data Access Mode | Simple database update and queries | Simple and complex queries with increasing use of data mining to recognise patterns in the data |
| Primary Database Model Employed | Relational and sometimes Hierarchical | Relational |
| Level of Detail | Detailed transactions | Often summarised |
| Historical Data | Current data only | Multiple years of data |
| Update Process | On-line, ongoing process as transactions are captured | Periodic process, once per week or once per month |
| Ease of Update | Routine and easy | Complex, must combine data from many sources-both internal and external |
| Data Integrity Issues | Each individual transaction must be closely edited | Major effort to “clean” and integrate data from multiple sources. |

OLTP

- On-line transaction processing (OLTP) is the traditional way of using a database
 - Short transactions (read/update few records) with ACID (**A**tomicity, **C**onsistency, **I**solation, **D**urability) properties
 - Normally, only the last version of data is stored in the database

OLTP vs OLAP



DSS & OLAP

- Decision support systems - help the executive, manager, analyst make faster and better decisions.
 - What were the sales volumes by region and product category for the last year?
 - Will a 10% discount increase sales volumes sufficiently?
- On-line analytical processing (OLAP) is an element of decision support systems (DSS)

Reasons for Building Data Warehouses

- Performance
 - OLAP applications need different organization of data
 - Complex OLAP queries would degrade OLTP performance
- Availability
 - Separation increases availability
 - Possibly the only way to query the disparate data sources


.... and the market is there!

Data Warehousing Tools Market

Worldwide Data Warehousing Tools Revenue by Vendor, 2003–2005



| | Revenue (\$M) | | | Share (%) | | | Growth (%) | |
|-------------------------------|---------------|---------|---------|-----------|------|------|---------------|---------------|
| | 2003 | 2004 | 2005 | 2003 | 2004 | 2005 | 2003– 2004 | 2004– 2005 |
| Oracle | 1,483.0 | 1,688.8 | 1,854.2 | 19.3 | 19.6 | 19.3 | 13.9 | 9.8 |
| IBM | 1,050.3 | 1,120.3 | 1,220.3 | 13.6 | 13.0 | 12.7 | 6.7 | 8.9 |
| SAS Institute | 826.1 | 922.6 | 1,021.6 | 10.7 | 10.7 | 10.7 | 11.7 | 10.7 |
| Microsoft | 630.5 | 802.4 | 985.3 | 8.2 | 9.3 | 10.3 | 27.3 | 22.8 |
| Business Objects | 367.1 | 404.1 | 450.0 | 4.8 | 4.7 | 4.7 | 10.1 | 11.4 |
| Teradata (division of NCR) | 325.4 | 390.0 | 423.0 | 4.2 | 4.5 | 4.4 | 19.8 | 8.5 |
| Cognos | 269.1 | 309.6 | 353.6 | 3.5 | 3.6 | 3.7 | 15.1 | 14.2 |
| Hyperion Solutions | 228.6 | 222.4 | 244.0 | 3.0 | 2.6 | 2.5 | -2.7 | 9.7 |
| Informatica | 154.2 | 168.3 | 211.6 | 2.0 | 2.0 | 2.2 | 9.1 | 25.8 |

Multidimensional Data



Dimensions: Product, Region, Date

Hierarchical summarization paths:






Operations

- Roll up: summarize data
- Drill down: go from higher level summary to lower level summary or detailed data
- Slice and dice: select and project
- Pivot: re-orient cube

Roll up


| Sales volume | | | |
|--------------|-------------|--------|--------|
| | Products | Store1 | Store2 |
| Q1 | Electronics | \$5,2 | \$5,6 |
| | Toys | \$1,9 | \$1,4 |
| | Clothing | \$2,3 | \$2,6 |
| | Cosmetics | \$1,1 | \$1,1 |
| Q2 | Electronics | \$8,9 | \$7,2 |
| | Toys | \$0,75 | \$0,4 |
| | Clothing | \$4,6 | \$4,6 |
| | Cosmetics | \$1,5 | \$0,5 |



| Sales volume | | | |
|--------------|-------------|--------|--------|
| | Products | Store1 | Store2 |
| Year 1996 | Electronics | \$14,1 | \$12,8 |
| | Toys | \$2,65 | \$1,8 |
| | Clothing | \$6,9 | \$7,2 |
| | Cosmetics | \$2,6 | \$1,6 |

Drill down


| Sales volume | | | |
|--------------|-------------|--------|--------|
| | Products | Store1 | Store2 |
| Q1 | Electronics | \$5,2 | \$5,6 |
| | Toys | \$1,9 | \$1,4 |
| | Clothing | \$2,3 | \$2,6 |
| | Cosmetics | \$1,1 | \$1,1 |
| Q2 | Electronics | \$8,9 | \$7,2 |
| | Toys | \$0,75 | \$0,4 |
| | Clothing | \$4,6 | \$4,6 |
| | Cosmetics | \$1,5 | \$0,5 |



| Sales volume | | | |
|--------------|-------------|--------|--------|
| | Electronics | Store1 | Store2 |
| Q1 | VCR | \$1,4 | \$1,4 |
| | Camcorder | \$0,6 | \$0,6 |
| | TV | \$2,0 | \$2,4 |
| | CD player | \$1,2 | \$1,2 |
| Q2 | VCR | \$2,4 | \$2,4 |
| | Camcorder | \$3,3 | \$1,3 |
| | TV | \$2,2 | \$2,5 |
| | CD player | \$1,0 | \$1,0 |

Slice and Dice


| Sales volume | | | |
|--------------|-------------|--------|--------|
| Q1 Q2 | Products | Store1 | Store2 |
| | Electronics | \$5,2 | \$5,6 |
| | Toys | \$1,9 | \$1,4 |
| | Clothing | \$2,3 | \$2,6 |
| | Cosmetics | \$1,1 | \$1,1 |
| | Electronics | \$8,9 | \$7,2 |
| | Toys | \$0,75 | \$0,4 |
| | Clothing | \$4,6 | \$4,6 |
| | Cosmetics | \$1,5 | \$0,5 |



| Sales volume | | |
|--------------|-------------|--------|
| Q1 Q2 | Products | Store1 |
| | Electronics | \$5,2 |
| | Toys | \$1,9 |
| | Electronics | \$8,9 |
| | Toys | \$0,75 |
| | | |

Pivot

| Sales volume | | | |
|--------------|-------------|--------|--------|
| | Products | Store1 | Store2 |
| | | | |
| Q1 | Electronics | \$5,2 | \$5,6 |
| | Toys | \$1,9 | \$1,4 |
| | Clothing | \$2,3 | \$2,6 |
| | Cosmetics | \$1,1 | \$1,1 |
| Q2 | Electronics | \$8,9 | \$7,2 |
| | Toys | \$0,75 | \$0,4 |
| | Clothing | \$4,6 | \$4,6 |
| | Cosmetics | \$1,5 | \$0,5 |

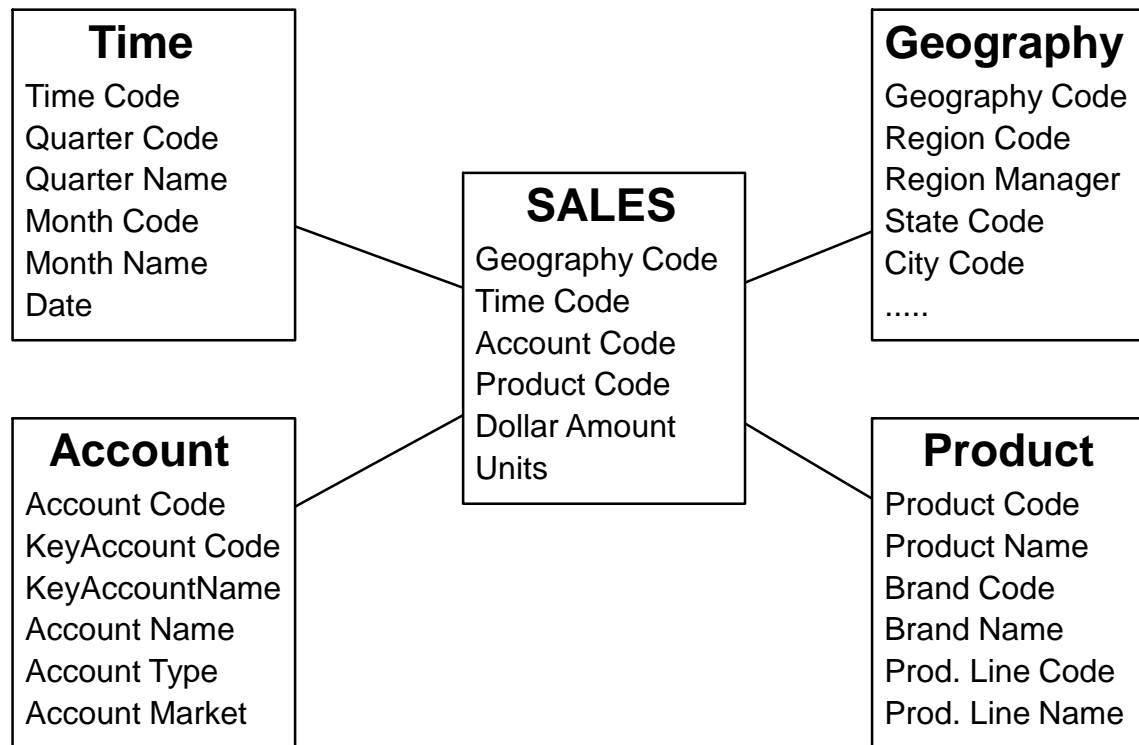


| Sales volume | | | |
|--------------|-------------|-------|--------|
| | Products | Q1 | Q2 |
| | | | |
| Store 1 | Electronics | \$5,2 | \$8,9 |
| | Toys | \$1,9 | \$0,75 |
| | Clothing | \$2,3 | \$4,6 |
| | Cosmetics | \$1,1 | \$1,5 |
| Store 2 | Electronics | \$5,6 | \$7,2 |
| | Toys | \$1,4 | \$0,4 |
| | Clothing | \$2,6 | \$4,6 |
| | Cosmetics | \$1,1 | \$0,5 |

1. Star Schema

- A star schema consists of one central **fact** table and several denormalized **dimension** tables.
- The **measures** of interest for OLAP are stored in the fact table (e.g. Dollar Amount, Units in the table SALES).
- For each dimension of the multidimensional model there exists a dimension table (e.g. Geography, Product, Time, Account) with all the **levels** of aggregation and the extra properties of these levels.

Star Schema...

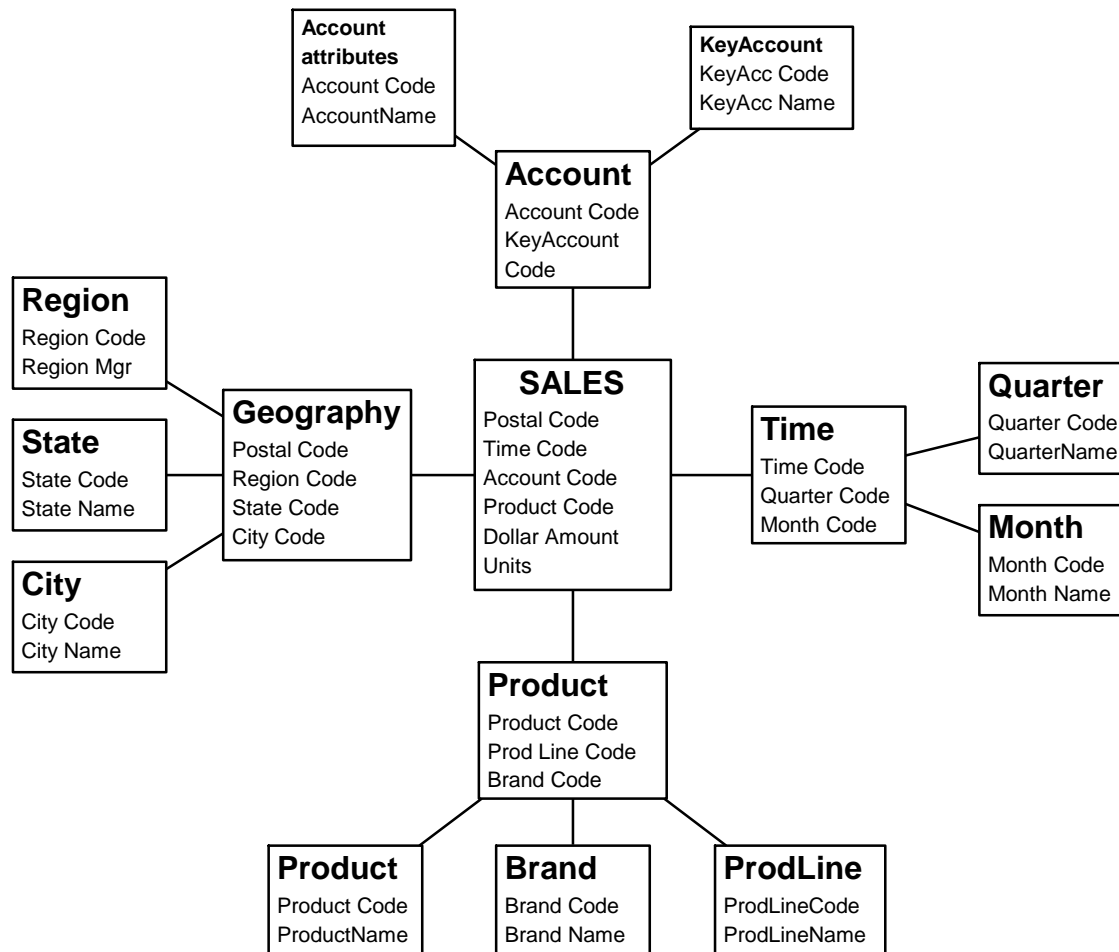




2. Snowflake Schema

- The normalized version of the star schema
- Explicit treatment of dimension hierarchies (each level has its own table)
- Easier to maintain, slower in query answering

Snowflake Schema...



Star Vs. Snowflake Schema

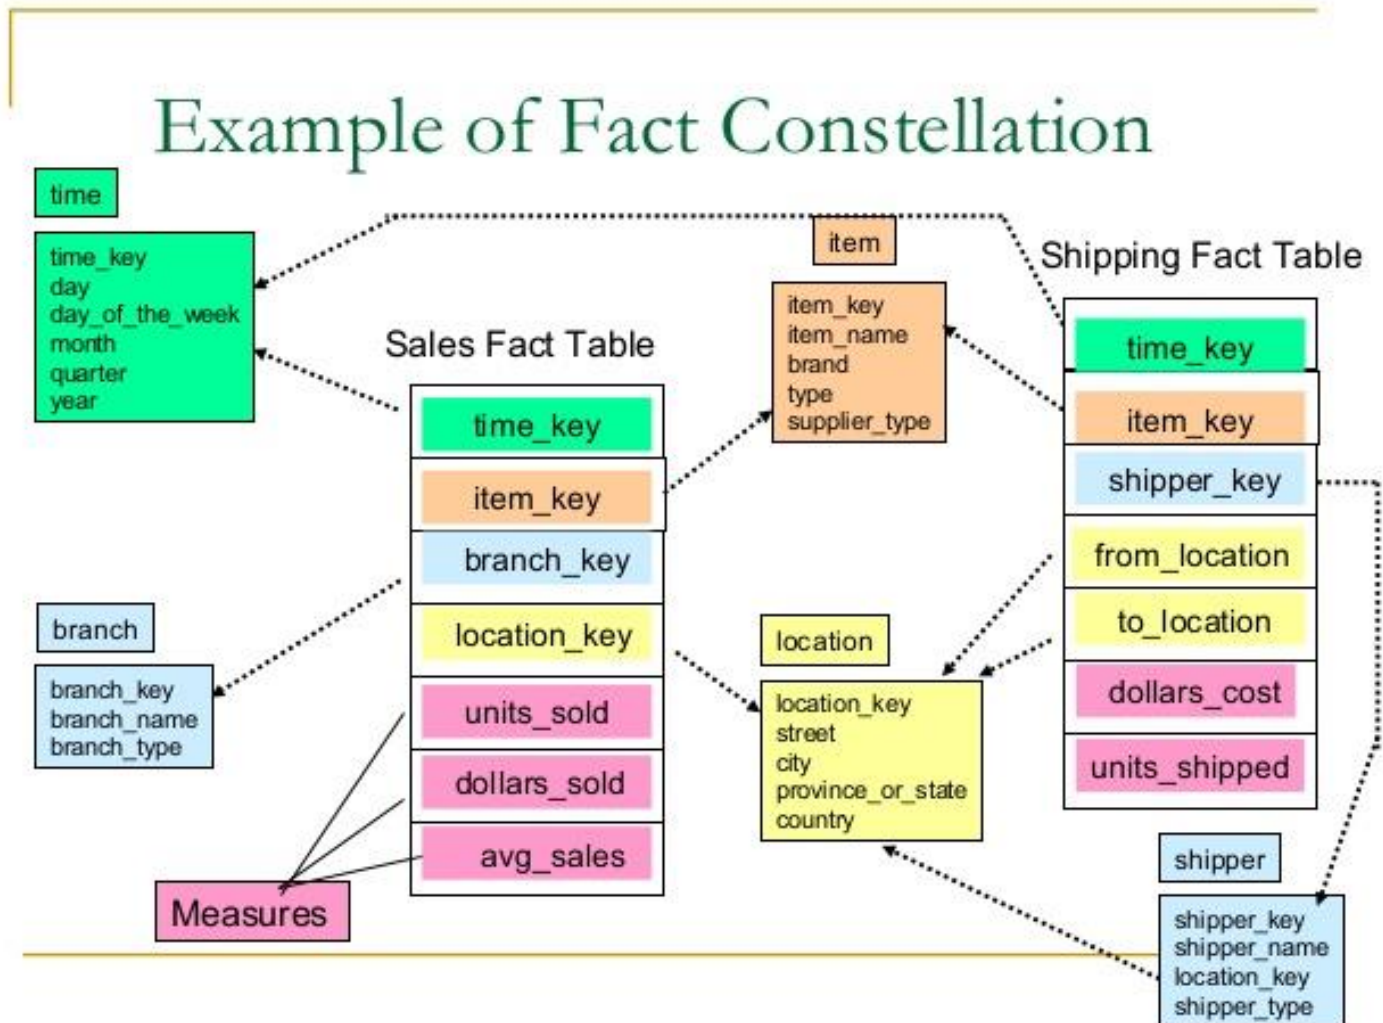
| | Snowflake Schema | Star Schema |
|-------------------------|---|--|
| Normalization | Can have normalized dimension tables. | Pure <u>denormalized</u> dimension tables. |
| Maintenance | Less redundancy so less maintenance. | More redundancy due to <u>denormalized</u> format so more maintenance. |
| Query | Complex Queries due to normalized dimension tables. | Simple queries due to pure <u>denormalized</u> design. |
| Joins | More joins due to normalization. | Less joins. |
| Usage guidelines | If you are concerned about integrity and duplication. | More than data integrity speed and performance is concern here. |



3. Fact Constellation

- Multiple fact tables that share many dimension tables
- Example: projected expense and the actual expense may share dimensional tables

Fact Constellation...





Operational Processes

- Propagate updates of source data to the warehouse
- Issues:
 - when to refresh
 - on every update
 - periodically
 - refresh policy set by administrator
 - how to refresh

Refreshment Techniques

- Full extract from base tables
- Incremental techniques
 - detect changes on base tables
 - data shipping (uses triggers to update warehouse tables)
 - transaction shipping (uses transaction log to ship transactions over to warehouse server for execution)

Accelerating the Refreshment Process

- Data Partitioning can be used to speed up refreshment
- Need to define a **partition key**





Partitioning


- Decomposing a large table or index into smaller and more manageable pieces.
- Data management operations can be performed at the partition level.
- Improves query performance.
- Can be implemented without requiring any modification to applications.

Partitioning Methods

List Partitioning


East Sales Region

New York
Virginia
Florida




West Sales Region

California
Oregon
Hawaii




Central Sales Region

Illinois
Texas
Missouri




Range Partitioning


January and
February




March and
April




May and
June



July and
August



Hash Partitioning



List Partitioning

- Enables control on how rows map to partitions
- Specify a list of discrete values for the partitioning key
- Example partitioning a sales table by region

```
CREATE TABLE sales_list (salesman_id NUMBER(5),  
    salesman_name VARCHAR2(30), sales_state VARCHAR2(20),  
    sales_amount NUMBER(10), sales_date DATE)  
PARTITION BY LIST(sales state)  
( PARTITION sales_west VALUES IN ('California',  
    'Hawaii'),  
    PARTITION sales_east VALUES IN ('New York', 'Virginia',  
    'Florida'),  
    PARTITION sales_central VALUES IN ('Texas', 'Illinois'),  
    )
```

Range Partitioning

- Maps data to partitions based on pages of partition key values.
- Most common type of partitioning often used with dates.

```
CREATE TABLE sales_range (salesman_id NUMBER(5),  
salesman_name VARCHAR2(30), sales_amount NUMBER(10),  
sales_date DATE) PARTITION BY RANGE(sales_date)  
(PARTITION sales_jan2002 VALUES LESS THAN  
(TO_DATE('02/01/2002', 'MM/DD/YYYY')),  
PARTITION sales_feb2002 VALUES LESS THAN  
(TO_DATE('03/01/2002', 'MM/DD/YYYY')),  
PARTITION sales_mar2002 VALUES LESS THAN  
(TO_DATE('04/01/2002', 'MM/DD/YYYY')),  
PARTITION sales_apr2002 VALUES LESS THAN  
(TO_DATE('05/01/2002', 'MM/DD/YYYY')), )
```

Hash Partitioning

- Used when range or list partitioning are not useful.
- Used when the amount of data to map into a given range is unknown.

```
CREATE TABLE sales_hash (salesman_id  
NUMBER(5), salesman_name VARCHAR2(30),  
sales_amount NUMBER(10), week_no NUMBER(2))  
PARTITION BY HASH(salesman_id)  
PARTITIONS 4  
STORE IN (data1, data2, data3, data4)
```




Indexes

- Bitmap Indexes
 - Reduced response time for large queries
 - Reduced storage requirements compared to other indexing techniques
 - Most often used in data warehouse applications
- B-tree Indexes
 - Most effective for high cardinality data
 - Used only for unique columns in data warehouse applications

Aggregated Tables

- In addition to base fact and dimension tables, data warehouse keeps aggregated (summary) data for efficiency.
- Two approaches
 - store as separate summary fact and dimension tables
 - add to the existing base tables

Aggregated Tables

- Separate sum-table

Sales table

| RID | City | Amount |
|-----|--------|--------|
| 1 | Athens | \$100 |
| 2 | N.Y. | \$300 |
| 3 | Rome | \$120 |
| 4 | Athens | \$250 |
| 5 | Rome | \$180 |
| 6 | Rome | \$65 |
| 7 | N.Y. | \$450 |

City-dimension sum table

| City | Amount |
|--------|--------|
| Athens | \$350 |
| N.Y. | \$750 |
| Rome | \$365 |

- Extend existing base tables

Extended Sales table

| RID | City | Amount | Level |
|-----|--------|--------|-------|
| 1 | Athens | \$100 | NULL |
| 2 | N.Y. | \$300 | NULL |
| 3 | Rome | \$120 | NULL |
| 4 | Athens | \$250 | NULL |
| 5 | Rome | \$180 | NULL |
| 6 | Rome | \$65 | NULL |
| 7 | N.Y. | \$450 | NULL |
| 8 | Athens | \$350 | City |
| 9 | N.Y. | \$750 | City |
| 10 | Rome | \$365 | City |

Materialized Views

- Summaries or Aggregate tables improve query execution times by pre-calculating expensive join and aggregation operations and storing the results in a table in the database.
- e.g. Create a table to contain the sums of sales by region and product.
- In Oracle you create a summary or aggregate table using a schema object called a materialized view.



Query Rewrite

- The query optimiser recognises when an existing materialized view can and should be used.
- The query is then transparently rewritten to use the materialized view.
- Rewriting queries to use materialized views instead of detail tables improves response time.



Materialized Views Behave Like Indexes

- The purpose is to increase query performance.
- Existence of a materialized view is transparent to SQL applications.
- They consume storage space.
- They must be updated when the underlying detail tables are modified.

Materialized View Example

```
CREATE MATERIALIZED VIEW product_sales_mv  
PCTFREE 0 TABLESPACE demo  
STORAGE (initial 8k next 8k pctincrease 0)  
BUILD IMMEDIATE  
REFRESH FAST  
ENABLE QUERY REWRITE  
AS SELECT p.prod_name, SUM(amount) AS  
dollar_sales, COUNT(*) AS cnt,  
COUNT(amount) AS cnt_amt  
FROM sales s, products p  
WHERE s.prod_id = p.prod_id  
GROUP BY prod_name;
```




References

- Elmasri, Navathe; *Fundamentals of Database Systems; 4th Ed.* Chapter 28.
- *Oracle 10g Data Warehousing Guide.*